

DDU-framework

January 27, 2026

```
[5]: %load_ext autoreload  
%autoreload 2
```

```
[7]: import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
from sklearn.linear_model import LinearRegression  
from sklearn.linear_model import Lasso, LassoCV  
import statsmodels.api as sm  
from sklearn.preprocessing import StandardScaler  
from datetime import datetime, timedelta  
from IPython.display import display, Math  
import re  
from arch import arch_model  
from arch.univariate import ZeroMean, GARCH, Normal  
from scipy.optimize import minimize  
from scipy.stats import norm  
import pickle  
from statsmodels.graphics.tsaplots import plot_acf  
import scipy.stats as stats
```

1 Functions

```
[9]: def plot_smoothed_residuals(date_strings, res, win, season):  
  
    # Calculate evenly spaced indices  
    indices = np.linspace(0, len(date_strings) - 1, num_labels_to_show, □  
    ↪dtype=int)  
    selected_dates = date_strings[indices]  
  
    # Take absolute value of residuals and apply rolling window smoothing □  
    ↪algorithm  
    y_res = abs(res).rolling(window=win).mean()  
  
    # Plot smoothed residuals  
    plt.plot(date_strings, y_res)
```

```

# Apply sampled ticks and rotated labels
ax = plt.gca()
ax.set_xticks(selected_dates)
ax.set_xticklabels(selected_dates, rotation=-45)
#ax.set_ylim([0, 1.75])

plt.title(season + " Season Smoothed Residuals")
plt.show()

return y_res

```

```

[10]: def garch_exog_loglik(params, y, x):
        """
        Log-likelihood for GARCH(1,1) with exogenous variable in variance equation
        """
        omega, alpha, beta, gamma = params
        n = len(y)

        # Initialize
        sigma2 = np.var(y) # Initial variance
        loglik = 0

        for t in range(1, n):
            # GARCH variance equation with exogenous variable
            sigma2 = omega + alpha * y[t-1]**2 + beta * sigma2 + gamma * x[t]

            # Ensure positive variance
            sigma2 = max(sigma2, 1e-8)

            # Log-likelihood contribution (assuming normal errors)
            loglik += -0.5 * (np.log(2 * np.pi) + np.log(sigma2) + y[t]**2 / sigma2)

        return -loglik # Return negative for minimization

```

```

[11]: def fit_garch_exog(y, x):
        """
        Fit GARCH(1,1) with exogenous variable in variance equation
        """
        # Initial parameter guesses
        init_params = [np.var(y) * 0.01, 0.1, 0.8, 0.01]

        # Parameter bounds and constraints
        bounds = [(1e-8, None),      # omega > 0
                  (0, 1),          # 0 <= alpha < 1
                  (0, 1),          # 0 <= beta < 1
                  (None, None)]    # gamma can be any value

```

```

constraints = [{'type': 'ineq',
                'funlambda params: 0.9999 - params[1] - params[2]}] # ↵
↳ alpha + beta < 1

# Explicit convergence criteria
options = {
    'ftolgtolxtolmaxiterdispTrue,          # Display convergence info
    'epsif hasattr(result, 'jac') else None
final_func_val = result.fun
final_params = result.x

# Calculate relative changes (approximate since we don't have
# second-to-last values)
# These are approximations - the actual algorithm tracks
# iteration-to-iteration changes
param_change_norm = np.linalg.norm(final_params - initial_params)
func_change = abs(final_func_val - initial_func_val)
relative_func_change = func_change / abs(initial_func_val) if
initial_func_val != 0 else func_change

# Print detailed convergence information
print(f"== Convergence Analysis ==")
print(f"Convergence achieved: {result.success}")
print(f"Termination message: {result.message}")
print(f"Number of iterations: {result.nit}")
print(f"Function evaluations: {result.nfev}")
print(f"")
print(f"== Tolerance Checks ==")

```

```

print(f"Final function value: {final_func_val}")
print(f"Function tolerance: {options['ftol']}")
print(f"Relative function change (total): {relative_func_change:.2e}")
print(f"")
print(f"Final gradient norm: {final_grad_norm}")
print(f"Gradient tolerance: {options['gtol']}")
print(f"Gradient criterion met: {final_grad_norm < options['gtol']} if"
     f" final_grad_norm else 'Unknown'")
print(f"")
print(f"Parameter change norm (total): {param_change_norm:.2e}")
print(f"Parameter tolerance: {options['xtol']}")
print(f"")
print(f"Final parameters: {final_params}")

return result

```

```

[12]: def forecast_volatility(result, y_train, x_test):
        """
        Simple volatility forecasting for GARCH-X model

        Parameters:
        -----
        result : fitted model result from fit_garch_exog
        y_train : training data (your y.values)
        x_test : test set exogenous variables (upstream inflow)

        Returns:
        -----
        volatility_forecast : array of forecasted volatilities
        """
        omega, alpha, beta, gamma = result.x

        # Convert to numpy array to avoid pandas indexing warnings
        x_test = np.array(x_test)

        # Get the last variance from training data
        sigma2 = np.var(y_train) # Initialize
        for t in range(1, len(y_train)):
            sigma2 = omega + alpha * y_train[t-1]**2 + beta * sigma2 + gamma * 0 #_
            # No x in training

        # Forecast volatilities
        forecasts = []
        last_y = y_train[-1]

        for t in range(len(x_test)):
            if t == 0:

```

```

        # First forecast uses last training observation
        sigma2 = omega + alpha * last_y**2 + beta * sigma2 + gamma * x_test[t]
    else:
        # Use previous forecast as input ( $E[y^2] = \sigma^2$ )
        sigma2 = omega + (alpha + beta) * sigma2 + gamma * x_test[t]

    forecasts.append(np.sqrt(sigma2)) # Convert to volatility

return np.array(forecasts)

```

```
[18]: def create_confidence_bands(y_test, y_pred_test, volatility_forecast,
                                confidence_level):
    """
    Create confidence intervals around predictions using GARCH volatility
    forecasts

    Parameters:
    -----
    y_test : array-like
        Actual test values (normalized residuals)
    y_pred_test : array-like
        Point predictions for test set
    volatility_forecast : array-like
        GARCH volatility forecasts
    confidence_level : float
        Confidence level (default 0.95 for 95% CI)

    Returns:
    -----
    dict : Dictionary with upper and lower bounds
    """
    from scipy.stats import norm

    # Calculate z-score for confidence level
    alpha = 1 - confidence_level
    z_score = norm.ppf(1 - alpha/2) # For 95% CI, z = 1.96

    # Create confidence bands around predictions
    lower_bound = y_pred_test - z_score * volatility_forecast
    upper_bound = y_pred_test + z_score * volatility_forecast

    return {
        'lower_bound': lower_bound,
        'upper_bound': upper_bound,
        'z_score': z_score
    }

```

```
[20]: def plot_forecast_with_uncertainty(y_test, y_pred, volatility_forecast):
    """
    Plot forecast with uncertainty bands in standard format

    Parameters:
    -----
    y_test : pandas Series or array
        Actual test values
    y_pred : pandas Series or array
        Point predictions for test set
    volatility_forecast : array
        GARCH volatility forecasts

    Returns:
    -----
    volatility_forecast : array (for consistency with your example)
    """
    # Create 95% confidence intervals
    upper = y_pred + 1.96 * volatility_forecast
    lower = y_pred - 1.96 * volatility_forecast

    # Plot
    plt.figure(figsize=(12, 2))
    plt.plot(y_test.index, y_test, label="Actual", color='blue')
    plt.plot(y_test.index, y_pred, label="Predicted", color='red')
    plt.fill_between(y_test.index, lower, upper, color='gray', alpha=0.3, label="95% CI")
    plt.xticks(rotation=-45)
    plt.legend()
    plt.title("Forecast with GARCH-X Uncertainty")
    plt.show()

    return volatility_forecast
```

```
[22]: def rescale(z, mean, std):
    return z*std + mean
```

2 Methods

3 Data Processing

3.1 Load Training and Test Data

```
[28]: # Load pickle file
with open("wet_ols_model.pkl", "rb") as f:
    wet_save_data = pickle.load(f)
```

```
# Unpack
wet_splits = wet_save_data["data"]
wet_y_pred = wet_save_data["y_pred"]
wet_residuals = wet_save_data["residuals"]
```

```
[30]: # Load pickle file
with open("dry_ols_model.pkl", "rb") as f:
    dry_save_data = pickle.load(f)

# Unpack
dry_splits = dry_save_data["data"]
dry_y_pred = dry_save_data["y_pred"]
dry_residuals = dry_save_data["residuals"]
```

4 Wet Season DDU Model

4.1 Calculate Date Indices

```
[34]: # Extract training dates
wet_date_strings = wet_splits["wet_X_train"].index.strftime('%Y-%m-%d').
    to_numpy()


# Set how many x-axis labels you want
num_labels_to_show = 15
```

4.2 Pre-Processing

```
[37]: # 1. Take absolute value of residuals

# 2. Apply a smoothing algorithm

wet_win = 5
wet_y_res = plot_smoothed_residuals(wet_date_strings, wet_residuals, wet_win, "Wet")
```



```
[39]: # 3. Normalize residuals
print(f"Before normalization - Mean: {wet_y_res.mean():.6f}, Std: {wet_y_res.
    ↪std():.6f}")


wet_y_res_norm = (wet_y_res - wet_y_res.mean()) / wet_y_res.std()
print(f"After normalization - Mean: {wet_y_res_norm.mean():.6f}, Std: {wet_y_res_norm.std():.6f}")
```

Before normalization - Mean: 0.068806, Std: 0.036420
After normalization - Mean: 0.000000, Std: 1.000000

```
[41]: # 4. Visualize transformed residuals


plt.hist(wet_y_res_norm, bins=50)
plt.title('Normalized & Smoothed Residual Distribution')
plt.show()
```

Normalized & Smoothed Residual Distribution



[43]: # 5. Recheck QQ Plots

```
stats.probplot(wet_y_res_norm, dist="norm", plot=plt)
plt.title("Q-Q Plot for Normalization Check")
plt.show()
```



4.3 Train ARCH-X Model

```
[46]: garch_x_col = 'outflow_m3hr_Lag_1'

[48]: # Drop any NaNs from lag alignment
garch_df = pd.concat([wet_splits["wet_X_train"][garch_x_col], wet_y_res_norm], axis=1).dropna()

X = garch_df[garch_x_col] # exogenous predictor (outflow lag 1)
y = garch_df[0]           # smoothed and normalized residuals

[50]: # Apply to your data

result = fit_garch_exog(y.values, X.values)

print("-----")
if result.success:
    omega, alpha, beta, gamma = result.x
    print("GARCH(1,1) with Exogenous Variable Results:")
    print(f"omega (constant): {omega:.6f}")
```

```

print(f"alpha (ARCH): {alpha:.6f}")
print(f"beta (GARCH): {beta:.6f}")
print(f"gamma (exogenous): {gamma:.6f}")
print(f"Log-likelihood: {-result.fun:.2f}")
else:
    print("Optimization failed:", result.message)

/var/folders/w1/xfz4_86j0zg0zz10g620bhz00000gn/T/ipykernel_74815/3796071886.py:3
2: OptimizeWarning: Unknown solver options: gtol, xtol
    result = minimize(garch_exog_loglik, init_params,
/opt/anaconda3/lib/python3.12/site-packages/scipy/optimize/_slsqp_py.py:437:
RuntimeWarning: Values in x were outside bounds during a minimize step, clipping
to bounds
    fx = wrapped_fun(x)

Optimization terminated successfully      (Exit mode 0)
    Current function value: 2416.296665515165
    Iterations: 24
    Function evaluations: 137
    Gradient evaluations: 24
==== Convergence Analysis ====
Convergence achieved: True
Termination message: Optimization terminated successfully
Number of iterations: 24
Function evaluations: 137

==== Tolerance Checks ====
Final function value: 2416.296665515165
Function tolerance: 1e-09
Relative function change (total): 1.76e-01

Final gradient norm: 32.826995880051065
Gradient tolerance: 1e-06
Gradient criterion met: False

Parameter change norm (total): 1.08e+00
Parameter tolerance: 1e-06

Final parameters: [0.17158209 0.81454127 0.          0.00832174]
-----
GARCH(1,1) with Exogenous Variable Results:
omega (constant): 0.171582
alpha (ARCH): 0.814541
beta (GARCH): 0.000000
gamma (exogenous): 0.008322
Log-likelihood: -2416.30

```


4.3.1 Forecast Volatility on Test Set

```
[53]: x_test = wet_splits["wet_X_test"][garch_x_col]
volatility_forecast = forecast_volatility(result, y.values, x_test)

[55]: y_pred_test = wet_y_pred
y_test = wet_splits['wet_y_test']

# Create confidence bands
confidence_level = 0.95
ci_bands = create_confidence_bands(y_test, y_pred_test, volatility_forecast,confidence_level)

# Plot everything together
rescaled_vol = rescale(volatility_forecast, wet_y_res.mean(), wet_y_res.std())
cond_vol = plot_forecast_with_uncertainty(y_test, y_pred_test, rescaled_vol)
```



4.3.2 Min/Max Estimated Std Analysis

```
[58]: # Raw Variance range on testing dataset

volatility_forecast.min(), volatility_forecast.max()
```

```
[58]: (0.4406710835041836, 0.9650427165790462)
```

```
[60]: # Rescaled variance range on testing dataset

rescale(volatility_forecast.min(), wet_y_res.mean(), wet_y_res.std()),
rescale(volatility_forecast.max(), wet_y_res.mean(), wet_y_res.std())
```

```
[60]: (0.08485540928014904, 0.10395310766759766)
```

5 Dry Season DDU Model

5.1 Calculate Date Indices


```
[64]: # Extract training dates
dry_date_strings = dry_splits["dry_X_train"].index.strftime('%Y-%m-%d').
    to_numpy()
```

5.2 Visualize Residuals

```
[67]: # 1. Take absolute value of residuals
# 2. Apply a smoothing algorithm

dry_win = 5

dry_y_res = plot_smoothed_residuals(dry_date_strings, dry_residuals, dry_win, "Dry")
```




```
[69]: # 3. Normalize residuals
print(f"Before normalization - Mean: {dry_y_res.mean():.6f}, Std: {dry_y_res.
    std():.6f}")

dry_y_res_norm = (dry_y_res - dry_y_res.mean()) / dry_y_res.std()
print(f"After normalization - Mean: {dry_y_res_norm.mean():.6f}, Std: {dry_y_res_norm.std():.6f}")
```

Before normalization - Mean: 0.210965, Std: 0.110301
After normalization - Mean: -0.000000, Std: 1.000000

```
[71]: # 4. Visualize transformed residuals


plt.hist(dry_y_res_norm, bins=50)
plt.title('Residual distribution')
plt.show()
```



```
[73]: # 5. Recheck QQ Plots

stats.probplot(dry_y_res_norm, dist="norm", plot=plt)
plt.title("Q-Q Plot for Normalization Check")
```

```
plt.show()
```



5.3 Train GARCH Model

```
[76]: # Drop any NaNs from lag alignment
garch_df = pd.concat([dry_splits["dry_X_train"][garch_x_col], dry_y_res_norm], axis=1).dropna()

X = garch_df[garch_x_col] # exogenous predictor (outflow lag 1)
y = garch_df[0]           # smoothed and normalized residuals
```

```
[78]: # Apply to your data
# result = fit_garch_exog(y.values, X.values)
result = fit_garch_exog(y.values, X.values)

if result.success:
    omega, alpha, beta, gamma = result.x
    print("GARCH(1,1) with Exogenous Variable Results:")
    print(f"omega (constant): {omega:.6f}")
    print(f"alpha (ARCH): {alpha:.6f}")
```

```

    print(f"beta (GARCH): {beta:.6f}")
    print(f"gamma (exogenous): {gamma:.6f}")
    print(f"Log-likelihood: {-result.fun:.2f}")
else:
    print("Optimization failed:", result.message)

/var/folders/w1/xfz4_86j0zg0zz10g620bhz0000gn/T/ipykernel_74815/3796071886.py:3
2: OptimizeWarning: Unknown solver options: gtol, xtol
    result = minimize(garch_exog_loglik, init_params,
/opt/anaconda3/lib/python3.12/site-packages/scipy/optimize/_slsqp_py.py:437:
RuntimeWarning: Values in x were outside bounds during a minimize step, clipping
to bounds
    fx = wrapped_fun(x)

Optimization terminated successfully      (Exit mode 0)
    Current function value: 2371.966948487659
    Iterations: 36
    Function evaluations: 221
    Gradient evaluations: 35
==== Convergence Analysis ====
Convergence achieved: True
Termination message: Optimization terminated successfully
Number of iterations: 36
Function evaluations: 221

==== Tolerance Checks ====
Final function value: 2371.966948487659
Function tolerance: 1e-09
Relative function change (total): 1.78e-01

Final gradient norm: 0.000813738502561852
Gradient tolerance: 1e-06
Gradient criterion met: False

Parameter change norm (total): 1.07e+00
Parameter tolerance: 1e-06

Final parameters: [0.17850484 0.79573642 0.0070291  0.04038951]
GARCH(1,1) with Exogenous Variable Results:
omega (constant): 0.178505
alpha (ARCH): 0.795736
beta (GARCH): 0.007029
gamma (exogenous): 0.040390
Log-likelihood: -2371.97

```


5.3.1 Forecast Volatility on Test Set

```
[81]: x_test = dry_splits["dry_X_test"][garch_x_col]
volatility_forecast = forecast_volatility(result, y.values, x_test)

[83]: y_pred_test = dry_y_pred
y_test = dry_splits['dry_y_test']

# Create confidence bands
ci_bands = create_confidence_bands(y_test, y_pred_test, volatility_forecast,confidence_level)

# Plot everything together
rescaled_vol = rescale(volatility_forecast, dry_y_res.mean(), dry_y_res.std())
cond_vol = plot_forecast_with_uncertainty(y_test, y_pred_test, rescaled_vol)
```



```
[85]: # Rescaled variance range on testing dataset

rescale(volatility_forecast.min(), dry_y_res.mean(), dry_y_res.std()),
rescale(volatility_forecast.max(), dry_y_res.mean(), dry_y_res.std())
```

```
[85]: (0.28393599397947666, 0.3289850736354223)
```

```
[ ]:
```