

# DIU-framework

January 27, 2026

```
[1]: %load_ext autoreload  
%autoreload 2
```

```
[3]: import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
from sklearn.linear_model import LinearRegression  
from sklearn.linear_model import Lasso, LassoCV  
import statsmodels.api as sm  
from sklearn.preprocessing import StandardScaler  
from datetime import datetime, timedelta  
from IPython.display import display, Math  
import re  
from arch import arch_model  
from arch.univariate import ZeroMean, GARCH, Normal  
import pickle  
import scipy.stats as stats  
from statsmodels.graphics.tsaplots import plot_acf
```

```
[4]: import preprocessing as preproc
```

```
[5]: def cfs_to_af(cfs_day):  
    """  
    Convert 'CFS-days' to acre-feet:  
  
    acre-ft = cfs_day / 1.9835  
    """  
    return cfs_day / 1.9835  
  
  
def af_to_m3(acft):  
    """Convert acre-feet to cubic metres (1 af = 1233.48 m³)."""  
    return acft * 1233.48  
  
def ft_to_m(ft):  
    return ft / 3.281
```

```

def kcfs_to_m3hr(kcfs):
    """(1000 cfs = 1kcfs = 28.32 m³/s --> m³/hr """
    return kcfs * 28.32 * 3600

```

```

[6]: def replace_outliers(data, max_val=1e8):
    """Replace values outside [-max_val, max_val] with average of nearest valid
    ↵neighbors"""
    arr = np.array(data, dtype=float)
    outlier_count = 0

    for i in range(len(arr)):
        if arr[i] > max_val or arr[i] < -max_val:
            outlier_count += 1
            # Find left neighbor
            left = None
            for j in range(i-1, -1, -1):
                if -max_val <= arr[j] <= max_val:
                    left = arr[j]
                    break

            # Find right neighbor
            right = None
            for j in range(i+1, len(arr)):
                if -max_val <= arr[j] <= max_val:
                    right = arr[j]
                    break

            # Replace with average of neighbors
            if left is not None and right is not None:
                arr[i] = (left + right) / 2
            elif left is not None:
                arr[i] = left
            elif right is not None:
                arr[i] = right
            else:
                arr[i] = 0

    print(f"Found and replaced {outlier_count} outliers")
    return arr

```

```

[11]: def timeseries_to_matrix(df):
    df['Date Time'] = pd.to_datetime(df['Date Time'])

    # Remove Feb 29
    df_no_leap = df[~((df['Date Time'].dt.month == 2) & (df['Date Time'].dt.day
    ↵== 29))].copy()

```

```

# Use original dayofyear but shift after Feb 29 removal
df_no_leap['Year'] = df_no_leap['Date Time'].dt.year
df_no_leap['Month'] = df_no_leap['Date Time'].dt.month
df_no_leap['Day'] = df_no_leap['Date Time'].dt.day

# Create day of year: Jan 1 = 1, Dec 31 = 365 (after removing Feb 29)
df_no_leap['DayOfYear'] = df_no_leap['Date Time'].dt.dayofyear

# Adjust for leap day removal: if after Feb 28, subtract 1 in leap years
leap_years = df_no_leap['Year'] % 4 == 0
after_feb28 = (df_no_leap['Month'] > 2)
df_no_leap.loc[leap_years & after_feb28, 'DayOfYear'] -= 1

matrix = df_no_leap.pivot_table(
    index='Year',
    columns='DayOfYear',
    values='Daily Inflow',
    aggfunc='first'
)

return matrix

```

```

[13]: def clean_variable_name(feature):
    '''Clean up variable names for LaTeX display'''
    # Replace inflow_m3hr_ with y
    if feature.startswith('inflow_m3hr_'):
        clean_name = feature.replace('inflow_m3hr_', 'y')
    # Replace outflow_m3hr_ with x
    elif feature.startswith('outflow_m3hr_'):
        clean_name = feature.replace('outflow_m3hr_', 'x')
    else:
        clean_name = feature

    # Handle Lag_n pattern - convert to subscript t-n
    lag_pattern = r'Lag_(\d+)'
    if re.search(lag_pattern, clean_name):
        # Extract the lag number
        lag_match = re.search(lag_pattern, clean_name)
        lag_num = lag_match.group(1)
        # Replace Lag_n with _{t-n}
        clean_name = re.sub(lag_pattern, f'_{{t-{lag_num}}}', clean_name)

    # Escape remaining underscores for LaTeX
    #clean_name = clean_name.replace('_', '\\_')

    return clean_name

```

# 1 Model Train Mode

```
[19]: framework = "DIU"
# framework = "DDU"
```

## 2 Data Load

```
[30]: bon_path = "/Users/elizacohn/Desktop/cascaded-hydro/streamflow-data-raw/
    ↴bonneville/"
tda_path = "/Users/elizacohn/Desktop/cascaded-hydro/streamflow-data-raw/dalles/"
```

### 2.1 Seasonal Analysis

```
[33]: bon_daily = pd.read_csv(bon_path + "bon-daily-inflow-calc.csv")
```

```
[35]: daily_inflow_col = 'BON.Flow-In.Ave.~1Day.1Day.CBT-REV [kcfs]'
bon_daily = bon_daily.rename(columns={daily_inflow_col: "Daily Inflow"})
```

```
[37]: bon_daily.head()
```

```
[37]:      Date Time Daily Inflow
0  31-Jul-1960 23:00      185.3
1  01-Aug-1960 23:00      170.9
2  02-Aug-1960 23:00      184.1
3  03-Aug-1960 23:00      168.9
4  04-Aug-1960 23:00      194.9
```

```
[39]: # Transform matrix in 365 x N years
```

```
inflow_matrix = timeseries_to_matrix(bon_daily)

print(f"Matrix shape: {inflow_matrix.shape}")
print(f"Years: {inflow_matrix.index.min()} to {inflow_matrix.index.max()}")
print(f"Days: {inflow_matrix.columns.min()} to {inflow_matrix.columns.max()}"
```

```
Matrix shape: (66, 365)
Years: 1960 to 2025
Days: 1 to 365
```

```
[41]: inflow_matrix.head()
```

```
[41]: DayOfYear      1      2      3      4      5      6      7      8      9      \
Year
1960        NaN     NaN     NaN     NaN     NaN     NaN     NaN     NaN     NaN
1961       106.7    98.2   97.4   98.2   99.3   124.2  110.8  122.4  134.8
1962       89.6    90.8  100.0   99.5   93.3   94.6  103.9  103.4  107.4
1963      127.2   130.0  140.1  139.9  137.5  142.4  128.0  128.0  135.4
```

```

1964      95.5  104.3   91.6   98.9   97.8  101.1   94.0  103.3  103.7
DayOfYear    10    ...   356   357   358   359   360   361   362   363  \
Year        ...
1960      NaN    ...  109.1   97.3   97.1   98.5   98.2  101.9   96.2   96.9
1961     123.7    ... 110.4  110.8  113.2  107.4   95.1   97.8   95.9   97.9
1962     144.7    ... 146.6  150.6  152.2  140.8  142.6  128.7  128.3  128.4
1963     133.0    ...  94.6   96.5   93.3   95.3   97.9  101.2   93.4   98.4
1964     103.9    ... 224.4  386.4  419.9  432.4  342.9    NaN  284.3  245.4

DayOfYear    364    365
Year
1960     103.7  100.0
1961      94.0   94.2
1962     136.5  139.1
1963      97.5   99.1
1964     190.1  186.9

[5 rows x 365 columns]

```

[43]: *## Calculate daily averages*

```

daily_averages = inflow_matrix.mean(axis=0, skipna=True)

# Create a single-row DataFrame for the averages
avg_row = pd.DataFrame([daily_averages], index=['Average'],
                       columns=inflow_matrix.columns)

# Concatenate original matrix with average row
inflow_matrix_with_avg = pd.concat([inflow_matrix, avg_row], axis=0)

```

[45]: *## Identify DoY with min and max flow*

```

# Get the average row
daily_averages = avg_row.loc['Average']

# Find day with minimum average inflow
min_day = daily_averages.idxmin()
min_value = daily_averages.min()

# Find day with maximum average inflow
max_day = daily_averages.idxmax()
max_value = daily_averages.max()

```

[47]: *def day\_to\_date(day\_of\_year):*  
 *# Create a date using pandas*

```

    date = pd.to_datetime("2023-01-01") + pd.Timedelta(days=int(day_of_year) - 1)
    return date.strftime("%B %d")

# Now this should work
print(f"Minimum inflow day: Day {min_day} ({day_to_date(min_day)}) = {min_value:.2f} kcfs")
print(f"Maximum inflow day: Day {max_day} ({day_to_date(max_day)}) = {max_value:.2f} kcfs")

```

Minimum inflow day: Day 275 (October 02) = 104.72 kcfs  
 Maximum inflow day: Day 157 (June 06) = 329.75 kcfs

[49]: *## Identify Seasons*

```

dry_start = (min_day - 45 - 1) % 365 + 1
dry_end = (min_day + 45 - 1) % 365 + 1
print(f"90-Day Dry season: Day {dry_start} ({day_to_date(dry_start)}) to Day {dry_end} ({day_to_date(dry_end)})")

wet_start = (max_day - 45 - 1) % 365 + 1
wet_end = (max_day + 45 - 1) % 365 + 1
print(f"90-Day Wet season: Day {wet_start} ({day_to_date(wet_start)}) to Day {wet_end} ({day_to_date(wet_end)})")

```

90-Day Dry season: Day 230 (August 18) to Day 320 (November 16)  
 90-Day Wet season: Day 112 (April 22) to Day 202 (July 21)

[51]: *# Create date labels for x-axis (using 2023 as reference year)*

```

def create_date_labels():
    dates = []
    date_labels = []
    base_date = datetime(2023, 1, 1)

    for day in range(1, 366): # Days 1-365
        current_date = base_date + timedelta(days=day - 1)
        dates.append(day)
        date_labels.append(current_date)

    return dates, date_labels

```

```
days, date_labels = create_date_labels()
```

```
# Alternative: More detailed plot with confidence intervals
def plot_with_percentiles():
```

```

month_starts = [1, 32, 60, 91, 121, 152, 182, 213, 244, 274, 305, 335] # ↴
↪Approximate month starts
month_labels = ['Jan', 'Feb', 'Mar', 'Apr', 'May', 'Jun',
                'Jul', 'Aug', 'Sep', 'Oct', 'Nov', 'Dec']

"""Plot with percentile bands showing variability"""

plt.figure(figsize=(15, 8))

# Calculate percentiles for each day (excluding the average row)
original_data = inflow_matrix_with_avg.iloc[:-1] # Exclude 'Average' row
p25 = original_data.quantile(0.25, axis=0)
p75 = original_data.quantile(0.75, axis=0)
p10 = original_data.quantile(0.10, axis=0)
p90 = original_data.quantile(0.90, axis=0)

# Plot average line
plt.plot(days, daily_averages.values, linewidth=3, color='darkblue', ↴
↪label='Average')

# Highlight extremes
plt.scatter([min_day], [min_value], color='red', s=100, zorder=5, ↴
↪label=f'Min avg: Day {min_day}')
plt.scatter([max_day], [max_value], color='green', s=100, zorder=5, ↴
↪label=f'Max avg: Day {max_day}')

# Add shaded regions for wet and dry seasons
if dry_start > dry_end: # Wraps around year
    plt.axvspan(dry_start, 365, alpha=0.2, color='orange', label=f'Dry ↴
↪Season (Day {dry_start}-365, 1-{dry_end})')
    plt.axvspan(1, dry_end, alpha=0.2, color='orange')
else:
    plt.axvspan(dry_start, dry_end, alpha=0.2, color='orange', label=f'Dry ↴
↪Season (Day {dry_start}-{dry_end})')

if wet_start > wet_end: # Wraps around year
    plt.axvspan(wet_start, 365, alpha=0.2, color='lightgreen', label=f'Wet ↴
↪Season (Day {wet_start}-365, 1-{wet_end})')
    plt.axvspan(1, wet_end, alpha=0.2, color='lightgreen')
else:
    plt.axvspan(wet_start, wet_end, alpha=0.2, color='lightgreen', ↴
↪label=f'Wet Season (Day {wet_start}-{wet_end})')

# Plot percentile bands
# plt.fill_between(days, p10.values, p90.values, alpha=0.2, ↴
↪color='lightblue', label='10th-90th percentile')

```


```

plt.fill_between(days, p25.values, p75.values, alpha=0.5, color='steelblue', label='25th-75th percentile')

plt.title('Bonneville Daily Inflow Patterns with Variability Bands (1960-2025)', fontsize=16, fontweight='bold')
plt.xlabel('Day of Year', fontsize=12)
plt.ylabel('Inflow (kcfs)', fontsize=12)
plt.grid(True, alpha=0.3)
plt.legend()
plt.xticks(month_starts, month_labels)
plt.tight_layout()
plt.show()

plot_with_percentiles()

```



## 2.2 AR Time Series Data Load

```
[54]: ## Load Bonneville Data
# Manually fixing timezone (from PST)

bon = pd.read_csv(bon_path + "bon-fullflow.csv")
# bon = pd.read_excel(bon_path + "bon-inflow-calc.xlsx")
bon['Date Time'] = pd.to_datetime(bon['Date Time'])
bon['Date Time'] = bon['Date Time'] + pd.Timedelta(hours=4)
```

```
[56]: bon.head()
```

```
[56]:          Date Time  BON.Flow-Gen.Ave.1Hour.1Hour.CBT-REV [kcfs] \
0 2018-01-01 00:00:00                      170.3
1 2018-01-01 01:00:00                      171.6
2 2018-01-01 02:00:00                      172.2
3 2018-01-01 03:00:00                      171.5
4 2018-01-01 04:00:00                      172.3

          BON.Flow-In.Inst.~6Hours.0.RFC-FCST [kcfs] \
0                               NaN
1                               NaN
2                               NaN
3                               NaN
4                               NaN

          BON.Flow-Out.Ave.1Hour.1Hour.CBT-REV [kcfs] \
0                      173.8
1                      175.1
2                      175.7
3                      175.0
4                      175.8

          BON.Flow-Spill.Ave.1Hour.1Hour.CBT-REV [kcfs]
0                      0.0
1                      0.0
2                      0.0
3                      0.0
4                      0.0
```

```
[58]: ## Load Dalles Data
# Manually fixing timezone (from PST)

tda = pd.read_csv(tda_path + "tda-fullflow.csv")
tda['Date Time'] = pd.to_datetime(tda['Date Time'])
tda['Date Time'] = tda['Date Time'] + pd.Timedelta(hours=4)
```

```
[60]: tda.head()
```

```
[60]:          Date Time  TDA.Flow-Gen.Ave.1Hour.1Hour.CBT-REV [kcfs] \
0 2018-01-01 00:00:00                      162.6
1 2018-01-01 01:00:00                      161.9
2 2018-01-01 02:00:00                      163.6
3 2018-01-01 03:00:00                      161.9
4 2018-01-01 04:00:00                      148.4

          TDA.Flow-In.Inst.~6Hours.0.RFC-FCST [kcfs] \
0                               NaN
1                               NaN
```

```

2                      NaN
3                      NaN
4                      NaN

TDA.Flow-Out.Ave.1Hour.1Hour.CBT-REV [kcfs] \
0                  163.7
1                  162.9
2                  164.7
3                  162.9
4                  149.4

TDA.Flow-Spill.Ave.1Hour.1Hour.CBT-REV [kcfs]
0                  0.0
1                  0.0
2                  0.0
3                  0.0
4                  0.0

```

## 2.3 Merge Datasets

```
[63]: inflow_col = "BON.Flow-In.Inst.~6Hours.0.RFC-FCST [kcfs]"
# inflow_col = "Calculated Hourly Inflow [kcfs]"
outflow_col = "TDA.Flow-Out.Ave.1Hour.1Hour.CBT-REV [kcfs]"
```

```
[65]: # Gather Bonneville 6-Hr inflow and convert to m3/hr
flow = bon[["Date Time", inflow_col]].dropna().rename(columns={inflow_col: "inflow_kcfs"})
flow['inflow_m3hr'] = flow['inflow_kcfs'].apply(lambda x: kcfs_to_m3hr(x))

# Merge in with Dalles total outflow data
flow = flow.merge(tda[["Date Time", outflow_col]], how='inner', on='Date Time').
    rename(columns={outflow_col: "outflow_kcfs"})
flow['outflow_m3hr'] = flow['outflow_kcfs'].apply(lambda x: kcfs_to_m3hr(x))

# Drop kcfs columns
flow = flow.drop(['inflow_kcfs', 'outflow_kcfs'], axis=1)
```

```
[67]: # flow['inflow_m3hr'] = replace_outliers(flow['inflow_m3hr'])
```

```
[69]: flow.head()
```

```
[69]:      Date Time    inflow_m3hr    outflow_m3hr
0 2018-01-05 02:00:00  1.881769e+07   18371750.4
1 2018-01-05 08:00:00  1.868770e+07   9940320.0
2 2018-01-05 14:00:00  1.891240e+07   21165235.2
3 2018-01-05 20:00:00  1.906584e+07   19656345.6
4 2018-01-06 02:00:00  1.879455e+07   19381075.2
```

## 2.4 Extract Seasonal Datasets

```
[72]: # Remove 2025 data (ncomplete seasons)
flow = flow[flow["Date Time"].dt.year < 2025]

# Create DoY column
doy = flow["Date Time"].dt.dayofyear

# Wet season mask
wet_mask = (doy >= wet_start) & (doy <= wet_end)

# Dry season mask
dry_mask = (doy >= dry_start) & (doy <= dry_end)

# Build the dataframes
wet_flow = flow[wet_mask].copy()
dry_flow = flow[dry_mask].copy()
```

```
[88]: # Visualize Flow Datasets
```

```
def plot_flow(df, tick_stride, season):
    plt.figure(figsize=(12, 4))

    # convert to strings for categorical x-axis
    x = df['Date Time'].dt.strftime("%Y-%m-%d")


    # plt.plot(x, df['outflow_m3hr'], label="TDA Outflow")
    plt.plot(x, df['inflow_m3hr'], label="BON Inflow")

    plt.title(f"{season} Season System Inflow", fontsize = 16)
    plt.ylabel('Flow [m³/hr]')
    plt.legend()


    # control tick density: show only every nth label
    plt.xticks(x[::tick_stride], rotation=45, ha='right')

    plt.tight_layout()
    plt.show()
```

```
[90]: plot_flow(wet_flow, 200, "Wet")
```



```
[92]: plot_flow(dry_flow, 200, "Dry")
```



### 3 Seasonal OLS Models

#### 3.1 Preprocessing

##### 3.1.1 Normalize Data (Z-Score)

```
[97]: ## Wet Season

wet_scaler = StandardScaler() # removes mean and converts to unit variance for
    ↪each column
wet_flow_norm = wet_scaler.fit_transform(wet_flow.copy().drop(['Date Time'], ↪
    ↪axis=1))
wet_flow_norm = pd.DataFrame(wet_flow_norm, columns=wet_flow.columns[1:])

[99]: print(f"{'Feature':<20} {'Mean':>12} {'Std Dev':>12} {'Variance':>12}")
print("-" * 58)
for i, col in enumerate(wet_flow.columns[1]):
```

```
print(f"{{col:<20} {wet_scaler.mean_[i]:>12.2e} {wet_scaler.scale_[i]:>12.
˓→2e} {wet_scaler.var_[i]:>12.2e}}")
```

Feature	Mean	Std Dev	Variance
<hr/>			
inflow_m3hr	2.58e+07	9.21e+06	8.49e+13
outflow_m3hr	2.37e+07	9.12e+06	8.32e+13

[101]: *## Dry Season*

```
dry_scaler = StandardScaler() # removes mean and converts to unit variance for
˓→each column
dry_flow_norm = dry_scaler.fit_transform(dry_flow.copy().drop(['Date Time'],
˓→axis=1))
dry_flow_norm = pd.DataFrame(dry_flow_norm, columns=dry_flow.columns[1:])
```

[103]: `print(f"{'Feature':<20} {'Mean':>12} {'Std Dev':>12} {'Variance':>12}")`

```
print("-" * 58)
for i, col in enumerate(dry_flow.columns[1:]):
    print(f"{{col:<20} {dry_scaler.mean_[i]:>12.2e} {dry_scaler.scale_[i]:>12.
˓→2e} {dry_scaler.var_[i]:>12.2e}}")
```

Feature	Mean	Std Dev	Variance
<hr/>			
inflow_m3hr	1.06e+07	2.47e+06	6.10e+12
outflow_m3hr	9.69e+06	2.79e+06	7.80e+12

### 3.1.2 Lag Features

[112]:

```
p = 4 + 1
up_feat = False
down_feat = True
x_col_name = 'outflow_m3hr'
y_col_name = 'inflow_m3hr'
```

[114]: *## Wet Season*

```
wet_flow_lag, feature_cols = preproc.create_lag_features_colname(wet_flow_norm.
˓→copy(), p, x_col_name, y_col_name, up_feat, down_feat)

# Add back in Date Time as index
wet_flow_lag.index = wet_flow["Date Time"]

# Drop rows with NaN values resulting from creating lag features
wet_flow_lag.dropna(inplace=True)
```

[116]: `wet_flow_lag.head()`

```
[116]:               inflow_m3hr  outflow_m3hr  inflow_m3hr_Lag_1 \
Date Time
2018-04-23 02:00:00      0.314129      0.284386      0.368393
2018-04-23 08:00:00      0.283256      0.083247      0.314129
2018-04-23 14:00:00      0.231272      0.494464      0.283256
2018-04-23 20:00:00      0.553028      0.544749      0.231272
2018-04-24 02:00:00      0.476324      0.516813      0.553028

                           inflow_m3hr_Lag_2  inflow_m3hr_Lag_3  inflow_m3hr_Lag_4
Date Time
2018-04-23 02:00:00      0.472174      0.428079      0.494493
2018-04-23 08:00:00      0.368393      0.472174      0.428079
2018-04-23 14:00:00      0.314129      0.368393      0.472174
2018-04-23 20:00:00      0.283256      0.314129      0.368393
2018-04-24 02:00:00      0.231272      0.283256      0.314129
```

```
[118]: ## Dry Season
```

```
dry_flow_lag, feature_cols = preproc.create_lag_features_colname(dry_flow_norm.
    ↪copy(), p, x_col_name, y_col_name, up_feat, down_feat)

# Add back in Date Time as index
dry_flow_lag.index = dry_flow["Date Time"]

# Drop rows with NaN values resulting from creating lag features
dry_flow_lag.dropna(inplace=True)
```

```
[120]: dry_flow_lag.head()
```

```
[120]:               inflow_m3hr  outflow_m3hr  inflow_m3hr_Lag_1 \
Date Time
2018-08-19 02:00:00      0.291111      0.182363      0.381574
2018-08-19 08:00:00      0.327404      0.182363      0.291111
2018-08-19 14:00:00      0.438675     -0.011127      0.327404
2018-08-19 20:00:00      0.606140     -0.003826      0.438675
2018-08-20 02:00:00      0.496479      0.029031      0.606140

                           inflow_m3hr_Lag_2  inflow_m3hr_Lag_3  inflow_m3hr_Lag_4
Date Time
2018-08-19 02:00:00      0.642226      0.591607      0.842928
2018-08-19 08:00:00      0.381574      0.642226      0.591607
2018-08-19 14:00:00      0.291111      0.381574      0.642226
2018-08-19 20:00:00      0.327404      0.291111      0.381574
2018-08-20 02:00:00      0.438675      0.327404      0.291111
```

### 3.2 Train / Test Split

```
[123]: train_start_year = '2018'
test_start_year = '2024'

[125]: ## Enforce Framework
if framework == "DDU":
    feature_cols = ['inflow_m3hr_Lag_1', 'outflow_m3hr_Lag_1']

if framework == "DIU":
    feature_cols = ['inflow_m3hr_Lag_1'] #, 'inflow_m3hr_Lag_2']

print(feature_cols)

['inflow_m3hr_Lag_1']

[127]: ## Split the data into training and testing sets

def split_train_test(flow_lag, feature_cols, y_col_name, test_start_year, season):
    # Train/test split
    train_data = flow_lag[flow_lag.index < test_start_year]
    test_data = flow_lag[flow_lag.index >= test_start_year]

    # Features + target
    X_train = train_data[feature_cols]
    y_train = train_data[y_col_name]
    X_test = test_data[feature_cols]
    y_test = test_data[y_col_name]

    # Return dictionary with season-prepended keys
    return {
        f"{season}_train_data": train_data,
        f"{season}_test_data": test_data,
        f"{season}_X_train": X_train,
        f"{season}_y_train": y_train,
        f"{season}_X_test": X_test,
        f"{season}_y_test": y_test,
    }

[129]: wet_splits = split_train_test(wet_flow_lag, feature_cols, y_col_name, test_start_year, "wet")
dry_splits = split_train_test(dry_flow_lag, feature_cols, y_col_name, test_start_year, "dry")

[131]: wet_splits["wet_X_train"].head()
```

```
[131]: inflow_m3hr_Lag_1
Date Time
2018-04-23 02:00:00      0.368393
2018-04-23 08:00:00      0.314129
2018-04-23 14:00:00      0.283256
2018-04-23 20:00:00      0.231272
2018-04-24 02:00:00      0.553028
```

### 3.3 Wet Season OLS Model

```
[134]: # Step 1: Add constant (intercept)
wet_X_train_const = sm.add_constant(wet_splits["wet_X_train"])

# Step 2: Fit OLS model
wet_model = sm.OLS(wet_splits["wet_y_train"], wet_X_train_const).
    fit(cov_type='HC1') # Heteroskedasticity-consistent

# Step 3: View summary
print(wet_model.summary())
```

OLS Regression Results

---

Dep. Variable:	inflow_m3hr	R-squared:	0.990
Model:	OLS	Adj. R-squared:	0.990
Method:	Least Squares	F-statistic:	1.979e+05
Date:	Tue, 27 Jan 2026	Prob (F-statistic):	0.00
Time:	11:34:19	Log-Likelihood:	1836.4
No. Observations:	2179	AIC:	-3669.
Df Residuals:	2177	BIC:	-3657.
Df Model:	1		
Covariance Type:	HC1		

---

	coef	std err	z	P> z	[0.025
0.975]					
-----					
const	-0.0004	0.002	-0.159	0.874	-0.005
0.004					
inflow_m3hr_Lag_1	0.9954	0.002	444.827	0.000	0.991
1.000					

---

Omnibus:	364.026	Durbin-Watson:	2.141
Prob(Omnibus):	0.000	Jarque-Bera (JB):	6683.016
Skew:	0.139	Prob(JB):	0.00
Kurtosis:	11.575	Cond. No.	1.12

---

Notes:

[1] Standard Errors are heteroscedasticity robust (HC1)

```
[136]: # Extract coefficients and p-values into a tidy DataFrame
coef_df = pd.DataFrame({
    'feature': wet_model.params.index,
    'coefficient': wet_model.params.values
    #'p_value': model.pvalues.values,
    #'t_value': model.tvalues.values
})

# Optional: sort by absolute coefficient size
coef_df = coef_df.reindex(coef_df['coefficient'].abs().
    sort_values(ascending=False).index)


print(coef_df)
```

	feature	coefficient
1	inflow_m3hr_Lag_1	0.995421
0	const	-0.000354

```
[138]: wet_residuals = wet_model.resid
wet_fitted_vals = wet_model.fittedvalues
```

```
[140]: # Checking for Heteroskedasticity
```

```
preproc.plot_variance("Wet Season OLS", wet_fitted_vals, wet_residuals,□
    bounds=False)
```



```
[142]: # Display Equation
```

```
params = wet_model.params
latex_eq = f"\hat{y} = {params.iloc[0]:.4f}"
for coef, feature in zip(params.iloc[1:], feature_cols):
    sign = "+" if coef >= 0 else "-"
    latex_eq += f" {sign} {feature} {coef:.4f}
```

```

coef_abs = abs(coef)
clean_feature = clean_variable_name(feature)
latex_eq += f" {sign} {coef_abs:.3f} {clean_feature}"

```

[144]: `display(Math(latex_eq))`

$$\hat{y} = -0.0004 + 0.995y_{t-1}$$

### 3.3.1 Evaluation


[147]: `wet_X_test_const = sm.add_constant(wet_splits["wet_X_test"]) # Add intercept term`  
`wet_X_test_const['const'] = np.ones(len(wet_splits["wet_X_test"]))`  
`wet_X_test_const = pd.DataFrame(wet_X_test_const, columns=wet_model.model.exog_names)`  
`wet_y_pred = wet_model.predict(wet_X_test_const)`

[149]: `# Dropping the first predicted value since not enough lag info`

```

preproc.plot_inflow_forecasts(wet_splits["wet_y_test"].iloc[1:], wet_y_pred.
    .iloc[1:])

```



[151]: `# Calculate model performance without the first predicted value`

```

preproc.print_test_stats(wet_splits["wet_y_test"].iloc[1:], wet_y_pred.iloc[1:])

```


RMSE: 0.105

MAE: 0.075

R<sup>2</sup>: 0.900

### 3.3.2 Residuals


```
[154]: plt.hist(wet_residuals, bins = 100)
plt.title('Wet Season Training Residuals')
plt.xlim([-5, .5])
plt.show()
```




```
[156]: mean = wet_residuals.mean()
std = wet_residuals.std()
(round(mean,3), round(std, 4))
```

```
[156]: (0.0, 0.1042)
```

```
[158]: plot_acf(wet_residuals)
plt.show()
```



```
[160]: stats.probplot(wet_residuals, dist="norm", plot=plt)
plt.title("Q-Q Plot for Normalization Check")
plt.show()
```



### 3.3.3 Data Export

```
[74]: if framework == "DDU":
    wet_data_to_save = {
        "data": wet_splits,
        "y_pred": wet_y_pred,
        "residuals": wet_residuals,
    }

    # Save to pickle
    with open("wet_ols_model.pkl", "wb") as f:
        pickle.dump(wet_data_to_save, f)

    print("Saved")
```

Saved

### 3.4 Dry Season OLS Model

```
[163]: # Step 1: Add constant (intercept)
dry_X_train_const = sm.add_constant(dry_splits["dry_X_train"])

# Step 2: Fit OLS model
dry_model = sm.OLS(dry_splits["dry_y_train"], dry_X_train_const).
    fit(cov_type='HC1') # Heteroskedasticity-consistent

# Step 3: View summary
print(dry_model.summary())
```

#### OLS Regression Results

```
=====
Dep. Variable: inflow_m3hr R-squared: 0.901
Model: OLS Adj. R-squared: 0.901
Method: Least Squares F-statistic: 1.645e+04
Date: Tue, 27 Jan 2026 Prob (F-statistic): 0.00
Time: 11:34:50 Log-Likelihood: -597.40
No. Observations: 2174 AIC: 1199.
Df Residuals: 2172 BIC: 1210.
Df Model: 1
Covariance Type: HC1
=====

=====

            coef      std err          z      P>|z|      [0.025
0.975]
-----
const      0.0025      0.007      0.371      0.711     -0.011
0.016
inflow_m3hr_Lag_1  0.9502      0.007     128.257      0.000      0.936
0.965
=====

Omnibus: 380.752 Durbin-Watson: 2.124
Prob(Omnibus): 0.000 Jarque-Bera (JB): 3037.768
Skew: 0.593 Prob(JB): 0.00
Kurtosis: 8.668 Cond. No. 1.03
=====
```

Notes:

[1] Standard Errors are heteroscedasticity robust (HC1)

```
[165]: # Extract coefficients and p-values into a tidy DataFrame
coef_df = pd.DataFrame({
    'feature': dry_model.params.index,
    'coefficient': dry_model.params.values
    #'p_value': model.pvalues.values,
```

```

    #'t_value': model.tvalues.values
})

# Optional: sort by absolute coefficient size
coef_df = coef_df.reindex(coef_df['coefficient'].abs().
                           sort_values(ascending=False).index)

print(coef_df)


```

	feature	coefficient
1	inflow_m3hr_Lag_1	0.950187
0	const	0.002509

[167]: dry\_residuals = dry\_model.resid  
dry\_fitted\_vals = dry\_model.fittedvalues

[169]: # Checking for Heteroskedasticity

```
preproc.plot_variance("LinReg", dry_fitted_vals, dry_residuals, bounds=False)
```



[171]: # Display Equation

```

params = dry_model.params
latex_eq = f"\hat{y} = {params.iloc[0]:.4f}"
for coef, feature in zip(params.iloc[1:], feature_cols):
    sign = "+" if coef >= 0 else "-"
    coef_abs = abs(coef)
    clean_feature = clean_variable_name(feature)
    latex_eq += f" {sign} {coef_abs:.3f} {clean_feature}"

```

[173]: display(Math(latex\_eq))


$$\hat{y} = 0.0025 + 0.950y_{t-1}$$

### 3.4.1 Evaluation

```
[176]: dry_X_test_const = sm.add_constant(dry_splits["dry_X_test"]) # Add intercept term
dry_X_test_const['const'] = np.ones(len(dry_splits["dry_X_test"]))
dry_X_test_const = pd.DataFrame(dry_X_test_const, columns=dry_model.model.exog_names)
dry_y_pred = wet_model.predict(dry_X_test_const)

[178]: # Dropping the first predicted value since not enough lag info

preproc.plot_inflow_forecasts(dry_splits["dry_y_test"].iloc[1:], dry_y_pred.iloc[1:])
```



```
[180]: # Calculate model performance without the first predicted value

preproc.print_test_stats(dry_splits["dry_y_test"].iloc[1:], dry_y_pred.iloc[1:])
```


RMSE: 0.322

MAE: 0.221

R<sup>2</sup>: 0.866

### 3.4.2 Residuals


```
[183]: plt.hist(dry_residuals, bins = 100)
plt.title('Dry Season Training Residuals')
plt.xlim([-1.5, 1.5])
plt.show()
```




```
[185]: mean = dry_residuals.mean()  
       std = dry_residuals.std()  
       (round(mean,3), round(std, 4))
```

```
[185]: (-0.0, 0.3186)
```

```
[187]: plot_acf(dry_residuals)  
       plt.show()
```



```
[189]: stats.probplot(dry_residuals, dist="norm", plot=plt)
plt.title("Q-Q Plot for Normalization Check")
plt.show()
```



```
[191]: # plt.scatter(dry_splits["dry_y_train"], dry_residuals)
# plot moving window mean and you will see a bit of structure
# a nonlinear model (nearest neighbor or neural net will go away)
```

```
[193]: if framework == "DDU":
    dry_data_to_save = {
        "data": dry_splits,
        "y_pred": dry_y_pred,
        "residuals": dry_residuals,
    }

    # Save to pickle
    with open("dry_ols_model.pkl", "wb") as f:
        pickle.dump(dry_data_to_save, f)

    print("Saved")
```

[ ]: